



# Oracle Database

Oracle Database

Gopal Mukkamala





# Contents

<b>1</b>	<b>Oracle RAC</b>	<b>5</b>
1.1	<b>Networks</b>	<b>5</b>
1.1.1	bonding	5
1.1.2	GNS	7
1.2	<b>Storage</b>	<b>7</b>
1.3	<b>Oracle</b>	<b>7</b>
1.3.1	RAC	7
1.4	<b>OCR</b>	<b>8</b>
1.5	<b>OLR</b>	<b>8</b>
1.6	<b>backups</b>	<b>9</b>
1.7	<b>Recovery</b>	<b>9</b>
1.8	<b>Single Instance to RAC</b>	<b>9</b>
1.9	<b>Patching</b>	<b>9</b>
1.10	<b>Performance</b>	<b>9</b>
1.11	<b>Upgrades</b>	<b>9</b>
1.12	<b>AddingNewNode</b>	<b>9</b>
1.13	<b>DeleteNode</b>	<b>9</b>
1.14	<b>AddClusterDBtoOEM</b>	<b>9</b>
1.15	<b>BuildDataGuard</b>	<b>9</b>
1.16	<b>Deinstall</b>	<b>9</b>
<b>2</b>	<b>Oracle Exadata</b>	<b>11</b>
2.1	<b>Networks</b>	<b>11</b>



# 1. Oracle RAC

## 1.1 Networks

A RAC environment is all about networks and the subsequent interconnect network traffic.

At a bare minimum we need 1 public, 1 VIP, 1 private & 3 SCAN IP addresses. The public and private networks should be on a different separate subnet. The SCAN IP should be on the same public network segment.

### 1.1.1 bonding

Network bonding is a way so that two interfaces can be simultaneously connected together for redundancy purposes.

Example:- eth0 ✓ eth1 ✓ The two interfaces will be bonded together and connected to a bond interface so that even if one physical network is disconnected the cluster can be still up and running.

```
Command 1.1.1 — HOME. cat /etc/sysconfig/network-scripts/ifcfg-bond0
DEVICE=bond0
NAME=bond0
TYPE=Bond
BONDING_MASTER=yes
IPADDR=172.16.232.18
PREFIX=24
ONBOOT=yes
BOOTPROTO=none
BONDING_OPTS="mode=1 miimon=100"
```

```
cat /etc/sysconfig/network-scripts/ifcfg-eth0
HWADDR=
```

```
TYPE="Ethernet"
BOOTPROTO="none"
DEFROUTE="yes"
PEERDNS="yes"
PEERROUTES="yes"
IPV4_FAILURE_FATAL="no"
IPV6INIT="yes"
IPV6_AUTOCONF="yes"
IPV6_DEFROUTE="yes"
IPV6_PEERDNS="yes"
IPV6_PEERROUTES="yes"
IPV6_FAILURE_FATAL="no"
NAME="eth0"
UUID=
\textbf{ONBOOT="yes"}
\textbf{MASTER=bond0}
\textbf{SLAVE=yes}

cat /etc/sysconfig/network-scripts/ifcfg-eth1
HWADDR=
TYPE=Ethernet
BOOTPROTO=none
DEFROUTE=yes
PEERDNS=yes
PEERROUTES=yes
IPV4_FAILURE_FATAL=no
IPV6INIT=yes
IPV6_AUTOCONF=yes
IPV6_DEFROUTE=yes
IPV6_PEERDNS=yes
IPV6_PEERROUTES=yes
IPV6_FAILURE_FATAL=no
NAME=eth1
UUID=
\textbf{ONBOOT=yes}
MASTER=bond0
SLAVE=yes
}

nmcli con reload
systemctl restart network

Verify
cat /proc/net/bonding/bond0
```

### 1.1.2 GNS

GNS VIP is a static IP which is configured in the DNS. DNS delegates queries to GNS which provides name resolution at that address.

```
Command 1.1.2 — GNS. srvctl add gns -i ip_address -d domain
srvctl config gns
srvctl relocate gns -n node1
srvctl start gns
srvctl status gns
srvctl stop gns
```

*ip\_address is the virtual ip of the gns.*

## 1.2 Storage

Data is stored in diskgroups at a minimum you should have the following diskgroups.

+DATA +FRA +RECO

The oracle database binaries will take up about 6.7 Gig of space.

```
Command 1.2.1 — oracleasm. fdisk /dev/sdb
mkfs.ext4 /dev/sdb1
oracleasm createdisk +DATA /dev/sdb1
fdisk /dev/sdc
mkfs.ext4 /dev/sdc1
oracleasm createdisk +FRA /dev/sdc1
fdisk /dev/sdd
mkfs.ext4 /dev/sdd1
oracleasm createdisk +FRA /dev/sdd1
oracleasm listdisks
DATA
FRA
RECO
```

## 1.3 Oracle

### 1.3.1 RAC

You can only have one installation of Oracle Clusterware running in a cluster. This means that the Oracle clusterware should be at the highest level of all the current running databases.

This means that Clusterware 12c will support databases running on 10,11 & 12.

There can be only one clusterware running on a given oracle home.

#### 1. Extended Distance Cluster

Extended distance clusters in a nutshell is a RAC implementation separated by Geographical distance. The distance should be approximately 50 Miles or so. The connection for the networks have to be in the mili second range.

In a extended distance cluster the parameter ASM\_PREFERRED\_READ\_GROUP is most notable as a node may benefit from reading from a local ASM copy rather than a Primary copy which could technically be in another data center.

In order to set the parameter you can issue the following command:-

In order to administer the GI home you have to set the environment ORACLE\_HOME to the grid home.

**Command 1.3.1 — HOME.** `ORACLE_HOME=/u01/app/12.1.0/grid; export ORACLE_HOME`

During installation, ownership of the entire path to the Grid home is changed to root. (/u01, /u01/app, /u01/app/12.1.0, /u01/app/12.1.0/grid)

**Command 1.3.2 — s.** `ALTER SYSTEM SET ASM_PREFERRED_READ_GROUP=DATA.SITE1 scope=spfile;`

The benefit of using Oracle extended RAC is extremely fast recovery times.

There are various key points to keep in mind when configuring extended clusters namely:-

- Spread the Workload Evenly Across the Sites in the Extended Cluster
- Add a Third Voting Disk to Host the Quorum Disk
- Configure the Nodes to Be Within the Proximity of a Metropolitan Area
- Use Host-Based Storage Mirroring with Oracle ASM Normal or High Redundancy

In an extended cluster it is advisable to add a third voting disk for recovery. The voting disk is utilized by the CSS process of a Oracle Clusterware.

The commands to add a third voting disk are as:-

**Command 1.3.3 — votingdisk.** Take a backup of existing voting disk.  
`$CRS_HOME/bin/ocrconfig -export /tmp/ocrbackup -s online`

List the current voting disks

`$crsctl query css votedisk`

Stop the CRS

`#crsctl stop crs`

`#crsctl check crs`

`#crsctl add css votedisk /voting_disk/vote_disk3 -force`

`#crsctl query css votedisk`

`#crsctl delete css votedisk /dev/raw/raw3 -force`

`#crsctl query css votedisk`

`#crsctl start crs`

## 1.4 OCR

OCR is Oracle cluster registry.

**Command 1.4.1 — OCR.** `#ocrconfig -add +DATA`  
`#ocrconfig -delete /ocrdata/ocr_1`  
`#ocrconfig -delete /ocrdata/ocr_2`

## 1.5 OLR

OLR is Oracle Local registry. OLR is a registry similar to OCR located on each node in a cluster, but contains information specific to each node.

**Command 1.5.1** `ocrcheck -local`  
`/etc/olr.loc`



## 1.6 backups

Managing backups in a RAC environment. Backups in a RAC environment are similar to a non-RAC environment.

```
Command 1.6.1 run { Allocate channel; }
```

## 1.7 Recovery

Managing recovery in a RAC environment. Only one node undertakes the responsibility of recovering in a RAC environment.

## 1.8 Single Instance to RAC

Converting a database from single instance to RAC is achieved either manually or through rconfig.

## 1.9 Patching

You can apply patches in a RAC environment by using the Oracle OPatch utility.

## 1.10 Performance

RAC wait events are usually identified by running AWR reports.

The typical RAC events are gc current block request gc cr block request gc current block 3-way gc current block busy gc cr block grant 2-way

The column CLUSTER\_WAIT\_TIME in V\$SQLAREA represents the wait time incurred by individual SQL statements for global cache events and will identify the SQL which may need to be tuned.

## 1.11 Upgrades

## 1.12 AddingNewNode

Adding new nodes in a RAC environment are archived by running the addnode.sh script.

## 1.13 DeleteNode

Adding new nodes in a RAC environment are archived by running the deletenode.sh script.

## 1.14 AddClusterDBtoOEM

Adding cluster database to OEM 12c.

## 1.15 BuildDataGuard

Building Dataguard in a RAC Environment.

## 1.16 Deinstall

```
Command 1.16.1 runInstaller -deinstall -h /u01/app/12.1.0/grid rm -rf /etc/oraInst.loc rm -rf /u01/app/12.1.0/ rm -rf /etc/oracle/
```





## 2. Oracle Exadata

### 2.1 Networks